Response-state ablation

Qwen3-VL-2B | TextVQA | 5% per layer, all layers

Baseline soft accuracy: 79.7%

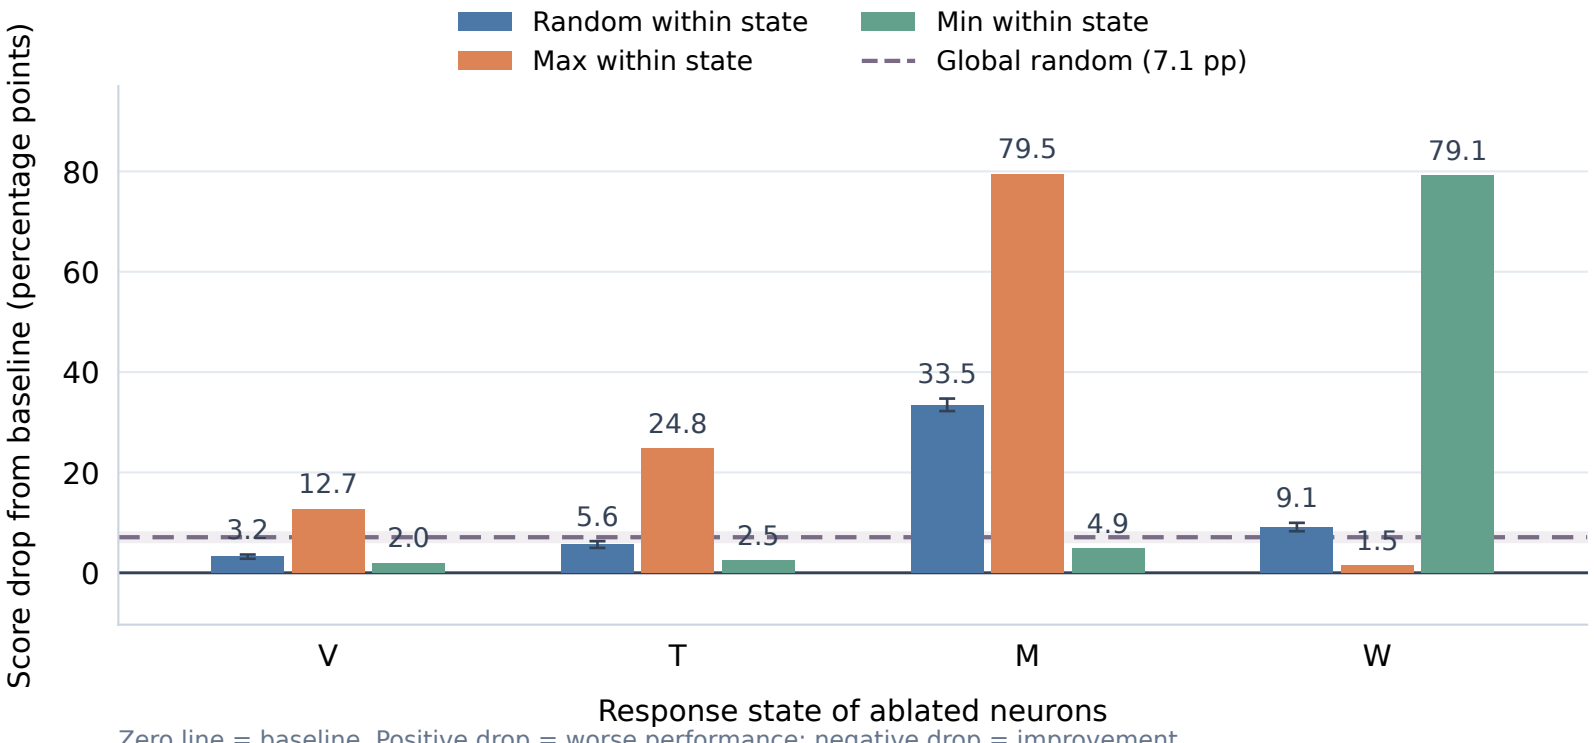

Zero line = baseline. Positive drop = worse performance; negative drop = improvement.

Error bars / band: +/- 1 SD across 10 random runs.