Sample-adaptive pruning: A1 vs A2

Qwen3-VL-2B | TextVQA | N = 1000 | Uniform ratio across layers per sample

Baseline soft accuracy: 79.7% | Measured results

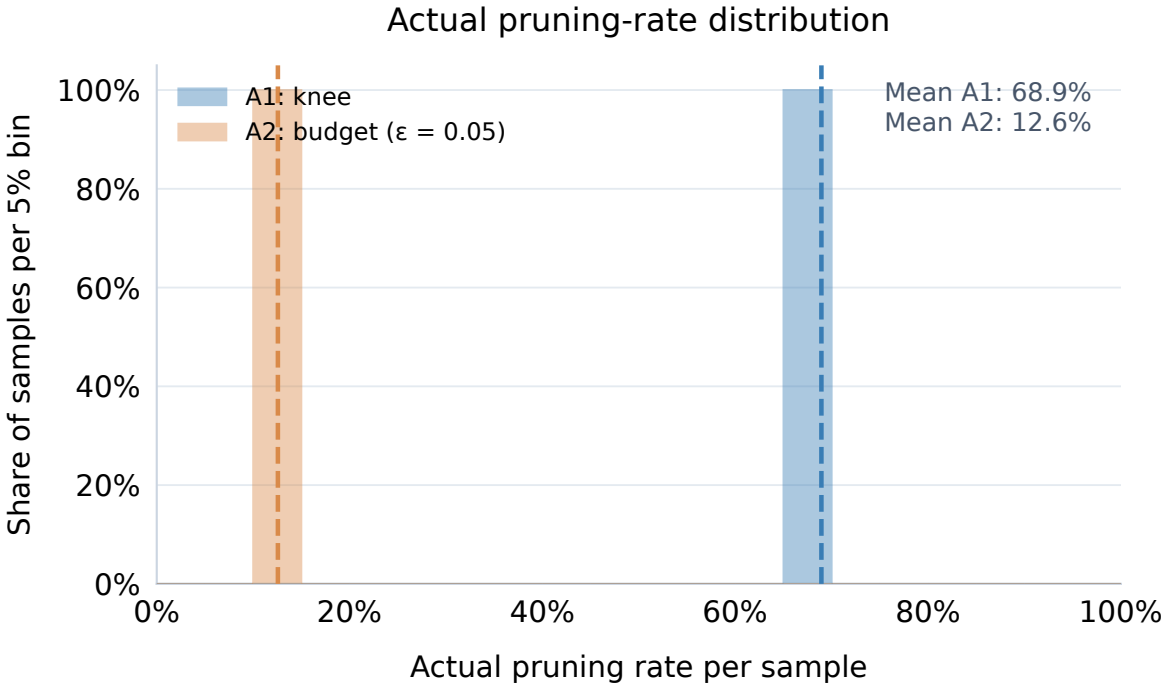

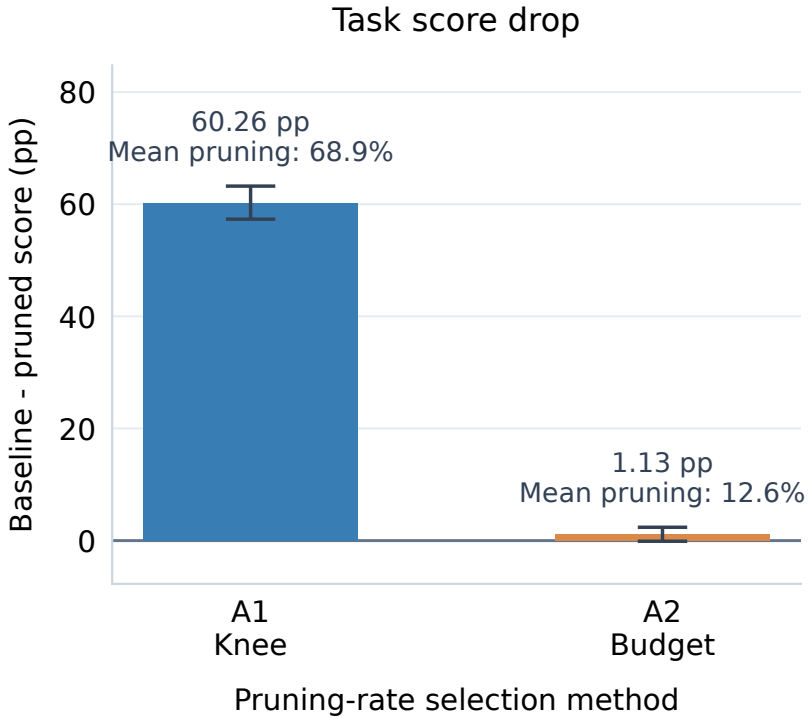

Dashed lines = mean pruning rates. Intervals = paired image/chart-group bootstrap 95% CIs.

Different removal amounts: compare quality and dose together. Positive drop = worse.