Task score vs response-score budget ε

Gemma-4-E2B-it | TextVQA | N = 1000

Measured results | Baseline soft accuracy: 65.0%

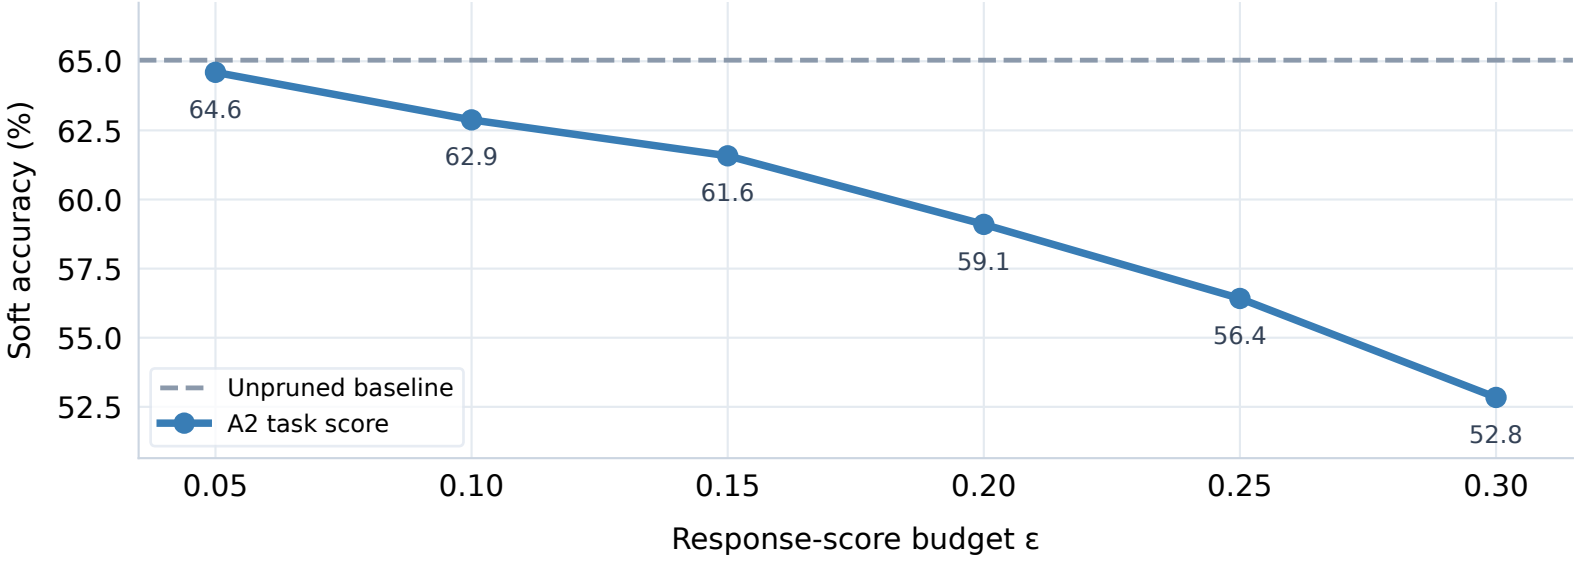

Mean actual pruning rate

ε=0.05	ε=0.10	ε=0.15	ε=0.20	ε=0.25	ε=0.30
11.8%	21.4%	30.0%	37.9%	45.2%	52.0%

Measured task scores; y-axis is zoomed. Confidence intervals remain available in summary.csv.