Task score vs response-score budget ε

LLaVA-1.5-7B | MMBench | N = 1000

Measured results | Baseline accuracy: 73.3%

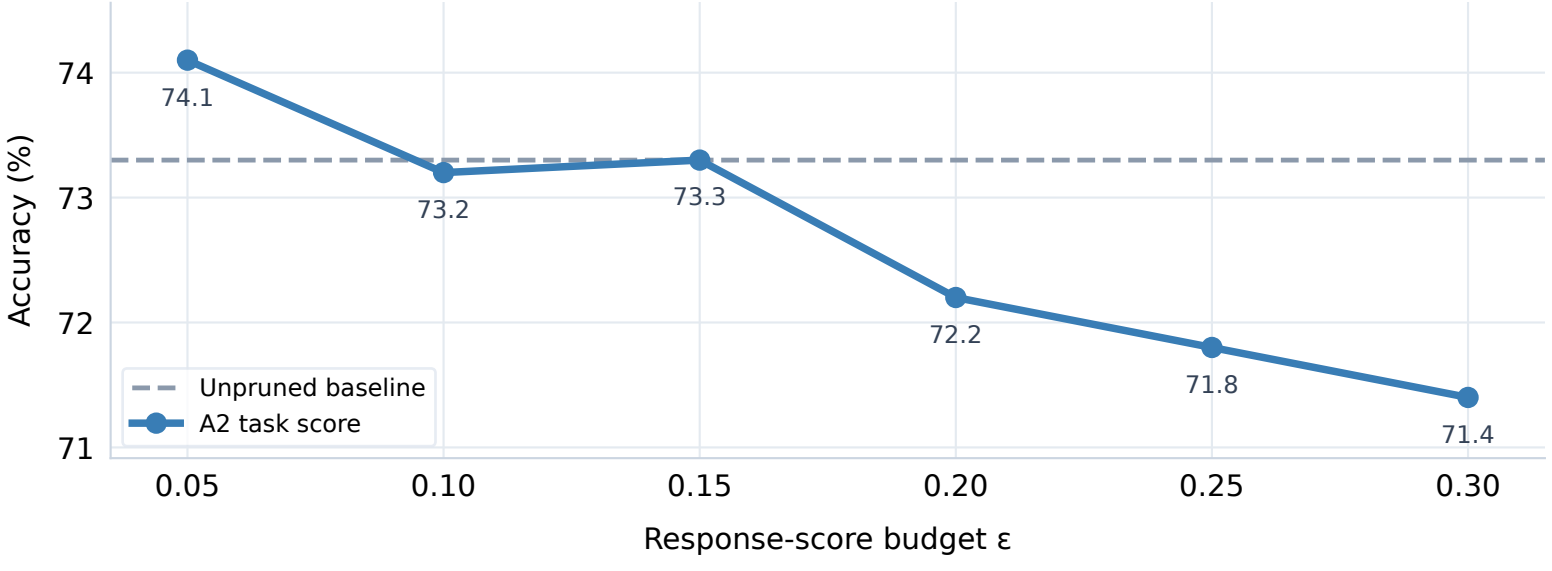

Mean actual pruning rate

ε=0.05	ε=0.10	ε=0.15	ε=0.20	ε=0.25	ε=0.30
9.4%	17.4%	24.7%	31.7%	38.2%	44.5%

Measured task scores; y-axis is zoomed. Confidence intervals remain available in summary.csv.