Task score vs response-score budget ε

Qwen3.5-2B | MMBench | N = 1000

Measured results | Baseline accuracy: 84.1%

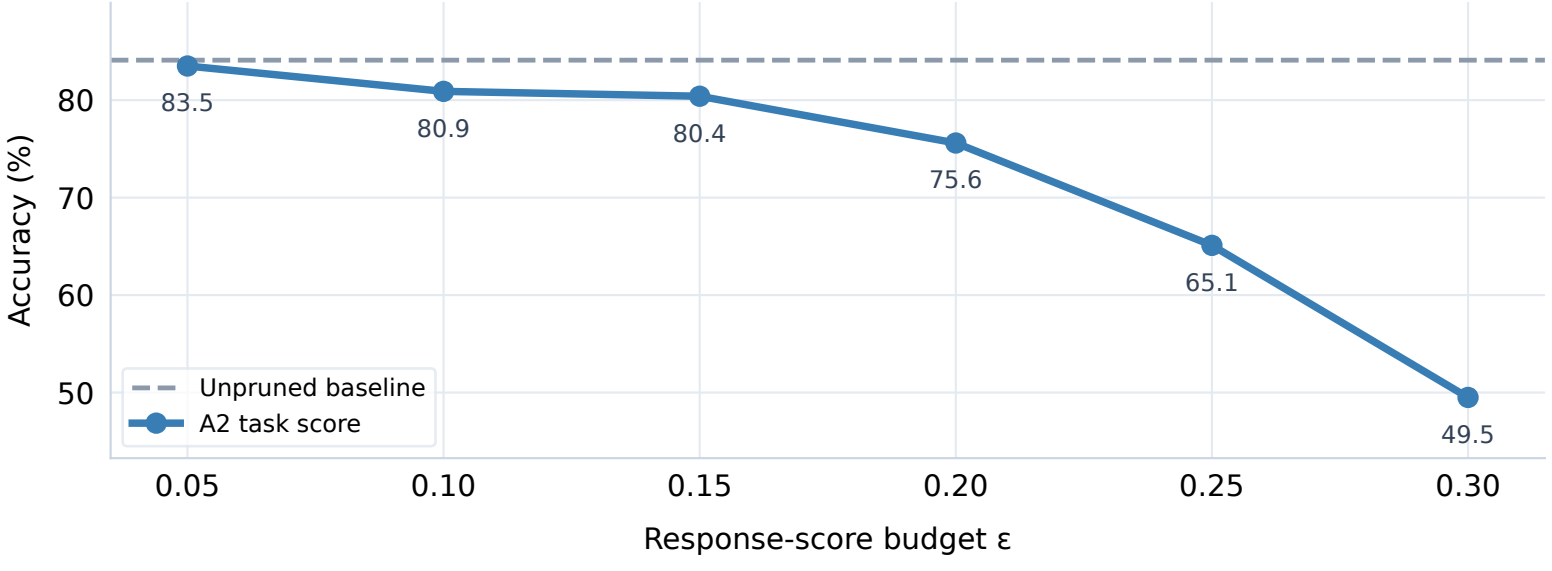

Mean actual pruning rate

ε=0.05	ε=0.10	ε=0.15	ε=0.20	ε=0.25	ε=0.30
8.9%	16.6%	23.6%	30.3%	36.7%	42.7%

Measured task scores; y-axis is zoomed. Confidence intervals remain available in summary.csv.