Task score vs response-score budget ε

Qwen3.5-2B | TextVQA | N = 1000

Measured results | Baseline soft accuracy: 77.6%

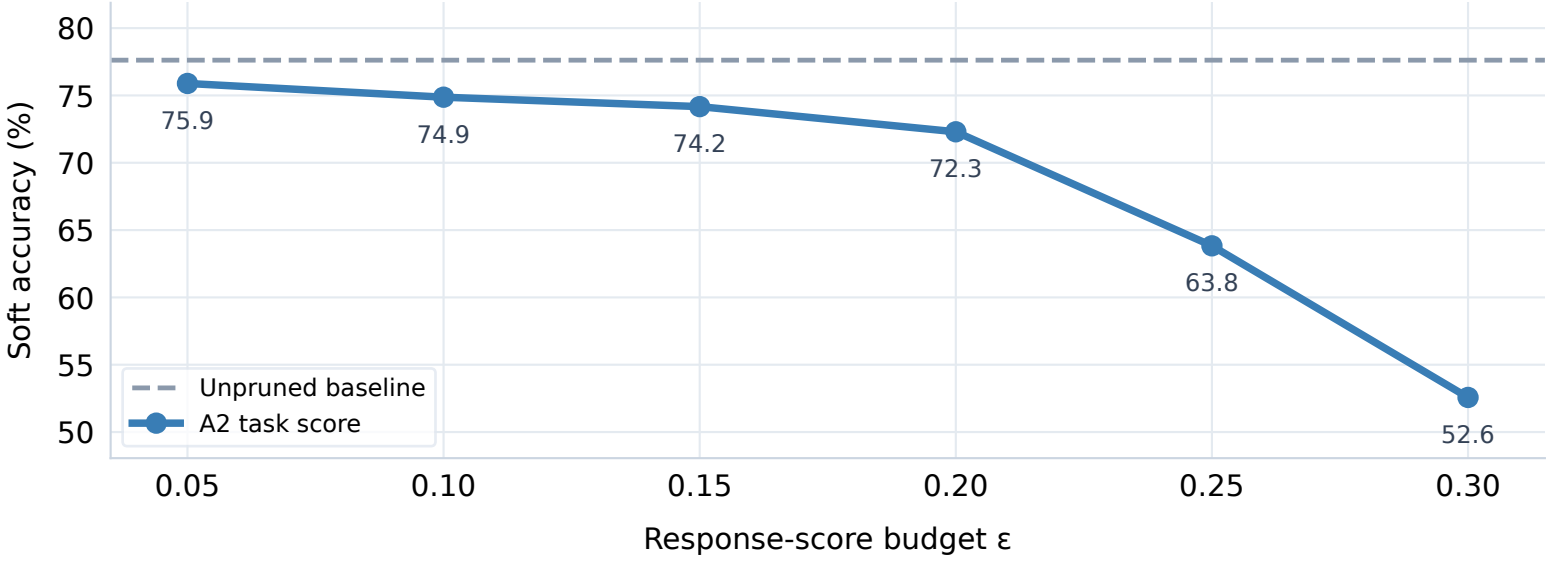

Mean actual pruning rate

ε=0.05	ε=0.10	ε=0.15	ε=0.20	ε=0.25	ε=0.30
9.0%	16.7%	23.8%	30.5%	36.9%	43.0%

Measured task scores; y-axis is zoomed. Confidence intervals remain available in summary.csv.