Task score vs response-score budget ε

Qwen3-VL-2B | TextVQA | N = 1000

Measured results | Baseline soft accuracy: 79.7%

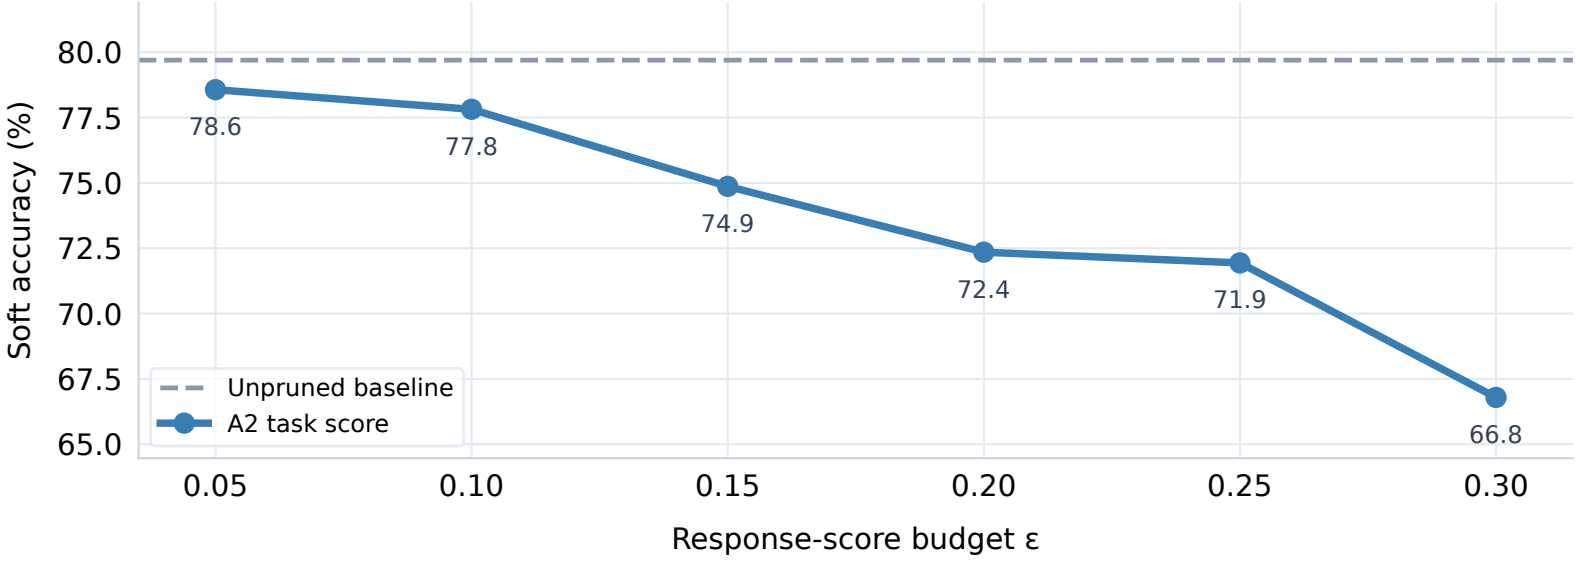

Mean actual pruning rate

ε=0.05	ε=0.10	ε=0.15	ε=0.20	ε=0.25	ε=0.30
12.6%	22.3%	31.0%	38.9%	46.2%	52.8%

Measured task scores; y-axis is zoomed. Confidence intervals remain available in summary.csv.