Uniform vs nonuniform pruning at the same response budget

Qwen3-VL-2B | MMBench | Sample-adaptive selection

SIMULATED EXAMPLE — not experimental results

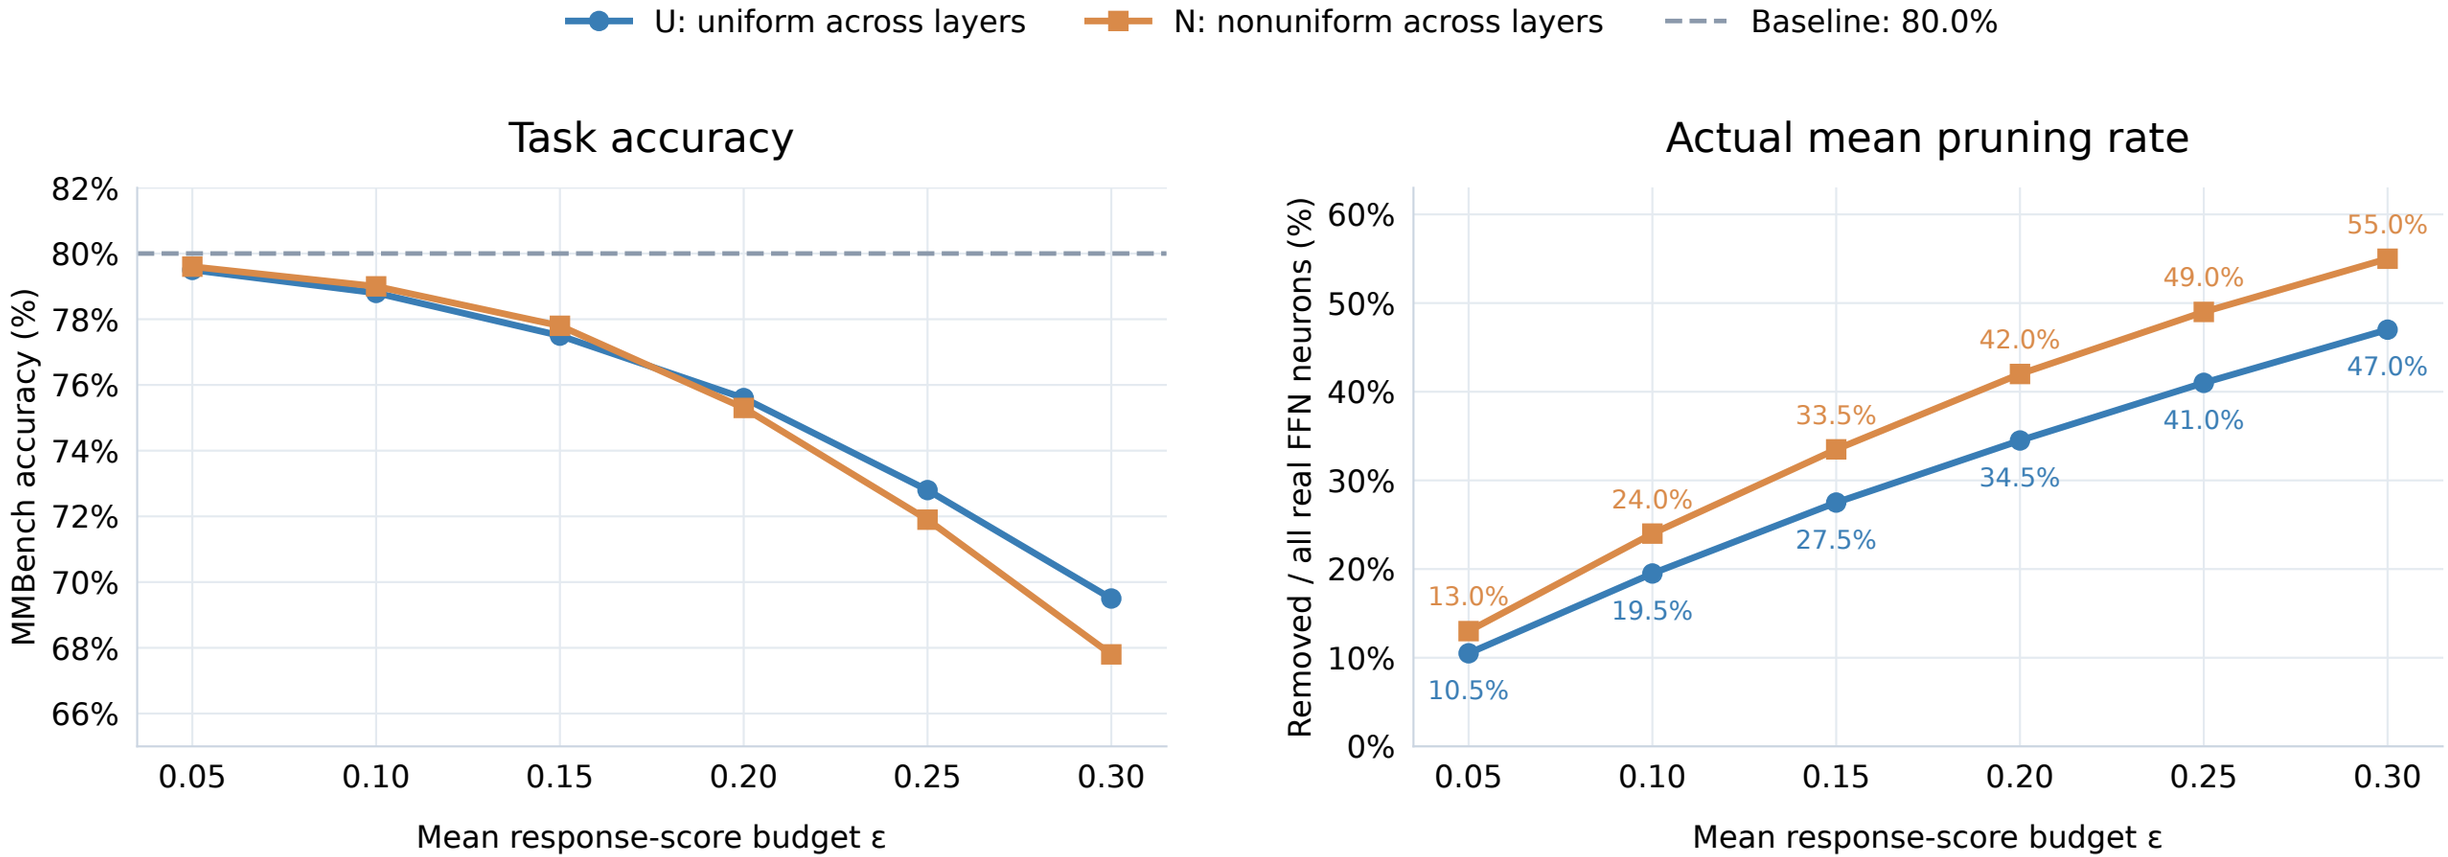

Actual mean response-score share removed (B)

Method	ε=0.05	ε=0.10	ε=0.15	ε=0.20	ε=0.25	ε=0.30
U: uniform	4.97%	9.96%	14.95%	19.95%	24.94%	29.94%
N: nonuniform	4.99%	9.99%	14.99%	19.99%	24.99%	29.99%

Budget = equal-layer mean response share. Pruning rate = total removed / total neurons.

All values are illustrative. Accuracy axis is zoomed. More removal does not guarantee better accuracy.