# Overall pruning: uniform vs nonuniform

4 models × 4 datasets | 1,000 frozen samples per run

Equal-weighted mean score drop from each run's own baseline

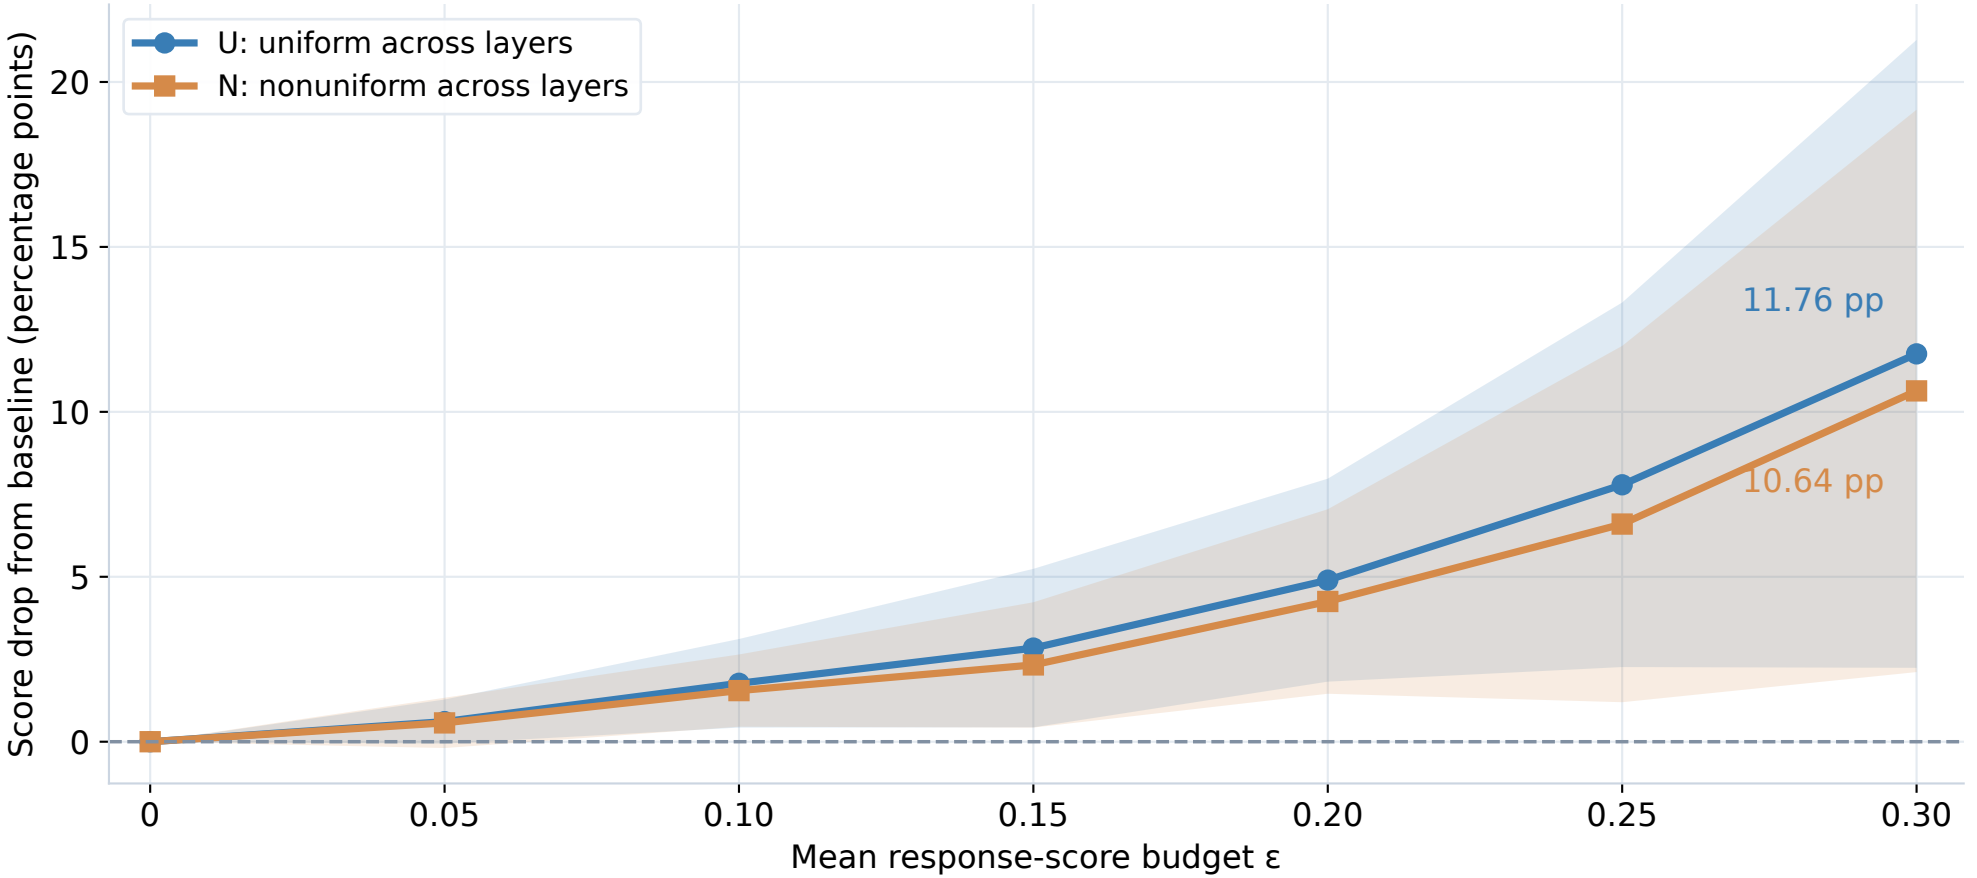


Shaded bands: ±1 SD across 16 runs (ddof=1); descriptive spread, not confidence intervals.

Positive drop = worse performance. Task metrics differ; each model–dataset run has equal weight.

Same response budget, not the same removal count. The point at ε=0 denotes the unpruned baseline.