Uniform vs nonuniform pruning at the same response budget

Qwen3.5-2B | TextVQA | Sample-adaptive selection

Measured results | N = 1000 paired samples per condition

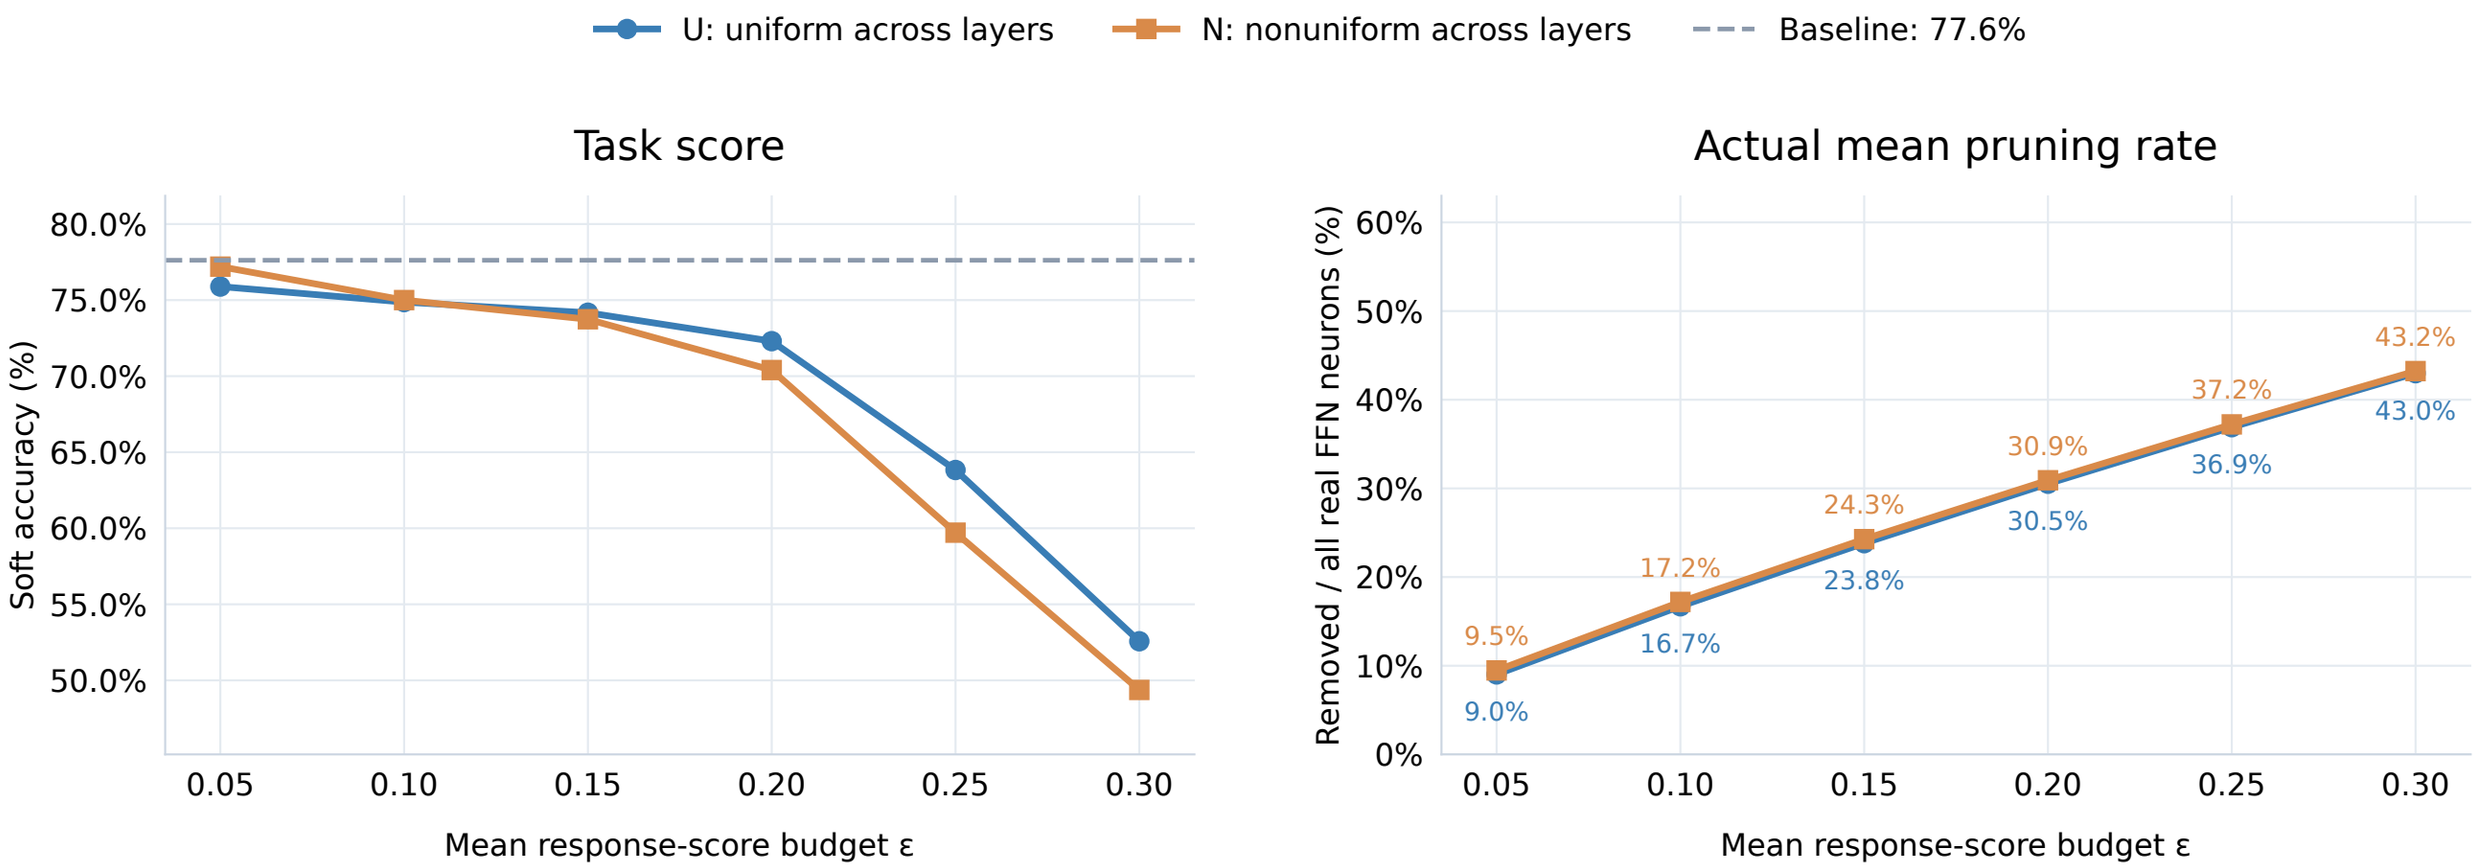

Actual mean response-score share removed (B)

Method	ε=0.05	ε=0.10	ε=0.15	ε=0.20	ε=0.25	ε=0.30
U: uniform	4.9949%	9.9943%	14.9940%	19.9937%	24.9935%	29.9930%
N: nonuniform	4.9998%	9.9998%	14.9997%	19.9997%	24.9997%	29.9997%

Budget = equal-layer mean response share. Pruning rate = total removed / total neurons.

Task-score axis is zoomed. More removal does not guarantee better accuracy.