Uniform vs nonuniform pruning at the same response budget

Qwen3-VL-2B | TextVQA | Sample-adaptive selection

Measured results | N = 1000 paired samples per condition

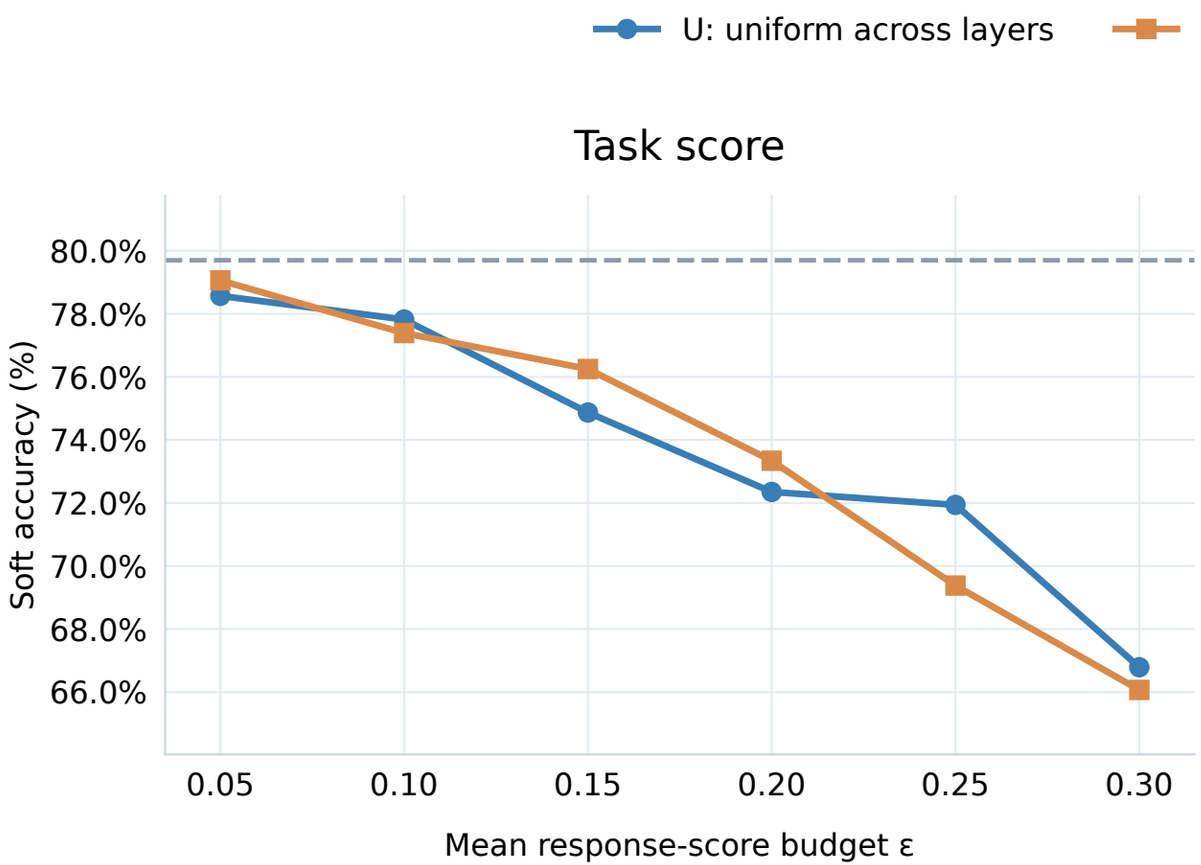

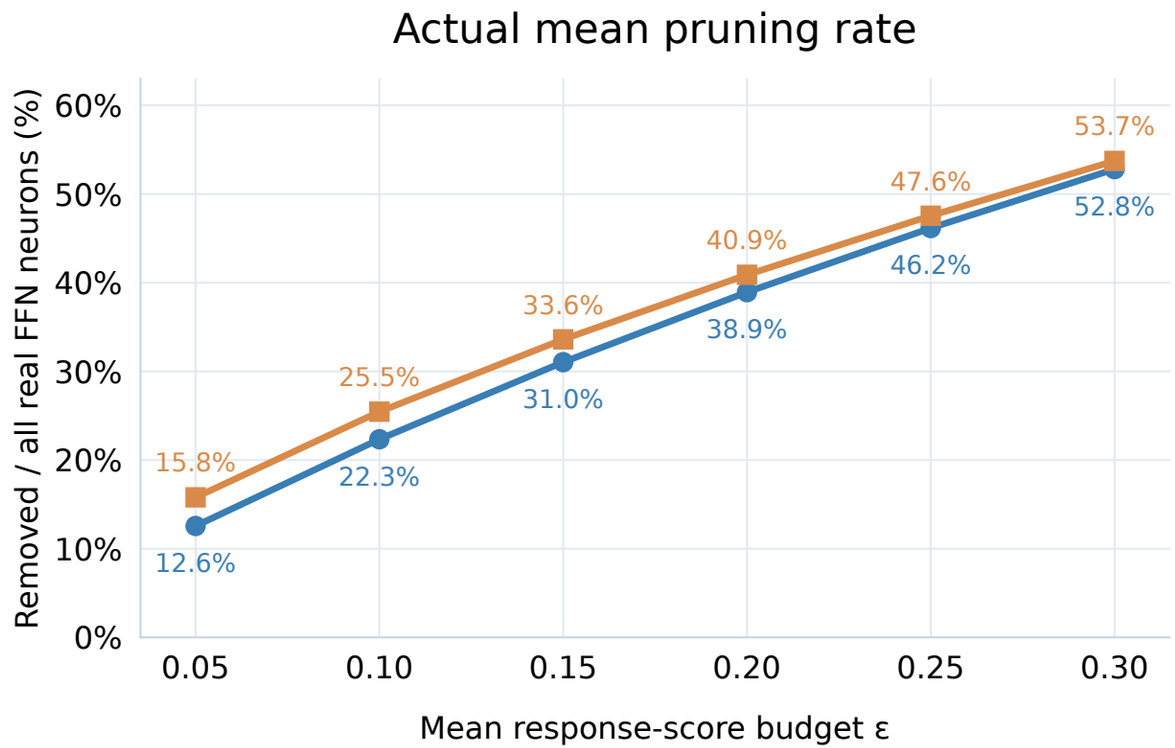

Actual mean response-score share removed (B)

Method	ε=0.05	ε=0.10	ε=0.15	ε=0.20	ε=0.25	ε=0.30
U: uniform	4.9961%	9.9956%	14.9950%	19.9946%	24.9941%	29.9936%
N: nonuniform	4.9999%	9.9998%	14.9998%	19.9998%	24.9998%	29.9998%

Budget = equal-layer mean response share. Pruning rate = total removed / total neurons.

Task-score axis is zoomed. More removal does not guarantee better accuracy.