Response-state ablation

LLaVA-1.5-7B | TextVQA | 1% per layer, all layers

Baseline soft accuracy: 48.64%

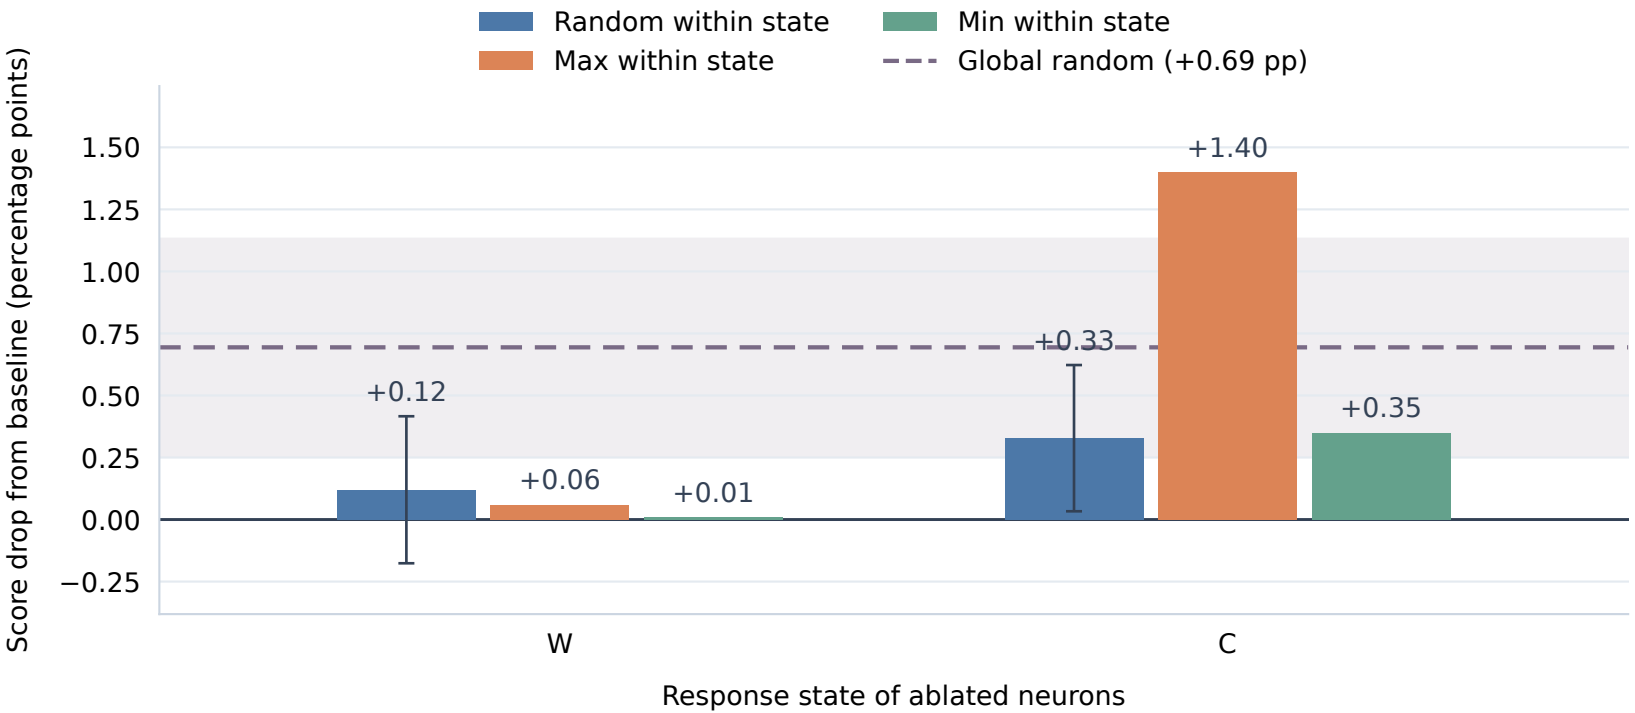

Zero line = baseline. Positive drop = worse performance; negative drop = improvement.

Error bars / band: +/- 1 SD across 10 random runs.

C neurons removed per sample: 3424.3/3520 target.