Response-state ablation

Qwen3-VL-2B | TextVQA | 1% per layer, all layers

Baseline soft accuracy: 79.70%

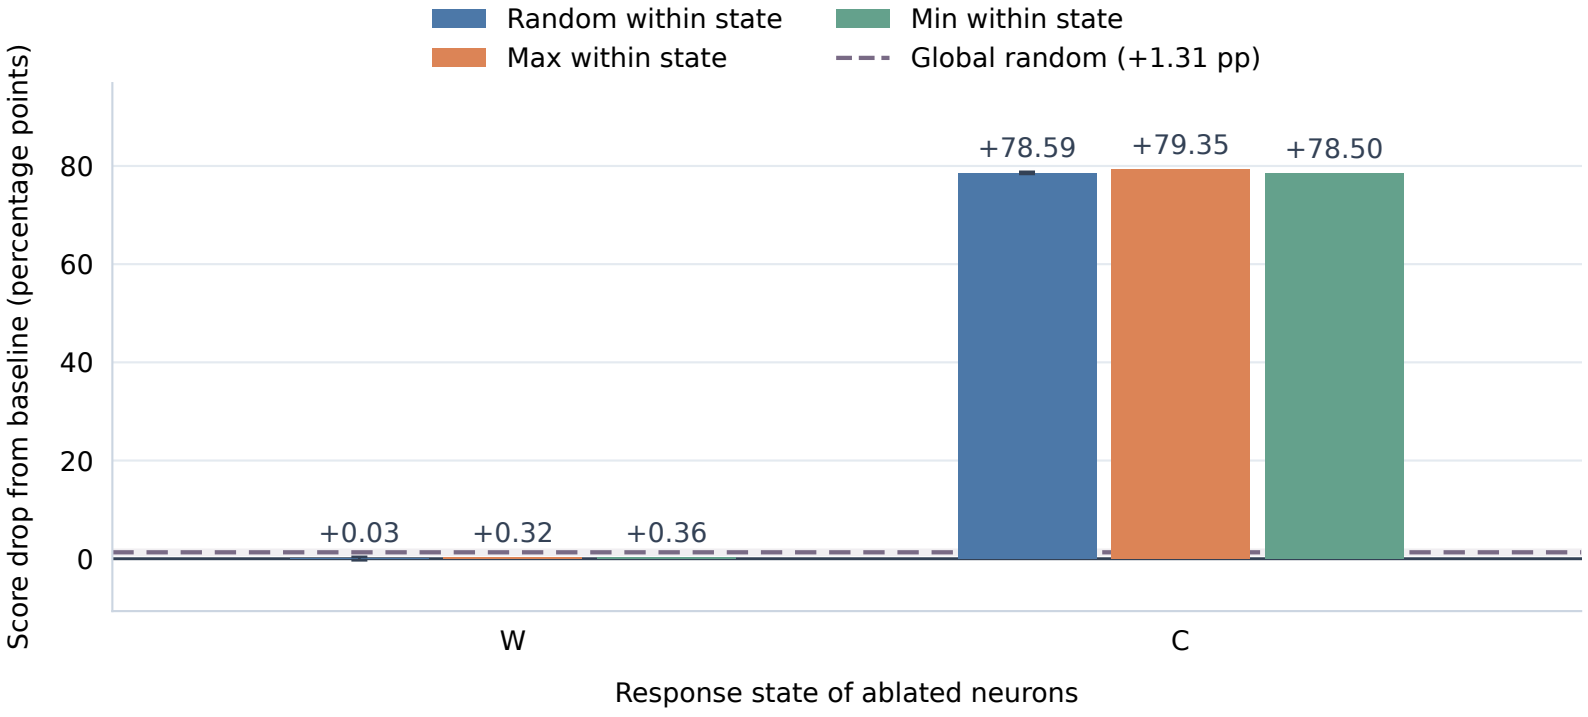

Zero line = baseline. Positive drop = worse performance; negative drop = improvement.

Error bars / band: +/- 1 SD across 10 random runs.

C neurons removed per sample: 1656.2/1708 target.