Inference throughput across pruning rates and batch sizes

Qwen3-VL-2B | ChartQA | 256 requests | MLP widths rounded up to multiples of 8

A40 | Measured GPU service throughput; median of 3 runs

Speedup uses original at the SAME batch size. First three panels share a req/s scale within each figure. N/A = comparison unavailable.

Right: blue = nonuniform faster; red = uniform faster. Actual aligned rates are in CSV. Natural EOS; max 32 new tokens; CPU preprocessing excluded.