Overall weak-neuron removal: A1 vs A2

4 models × 4 datasets | 1,000 frozen samples per run

Equal-weighted mean score drop from each run's own baseline

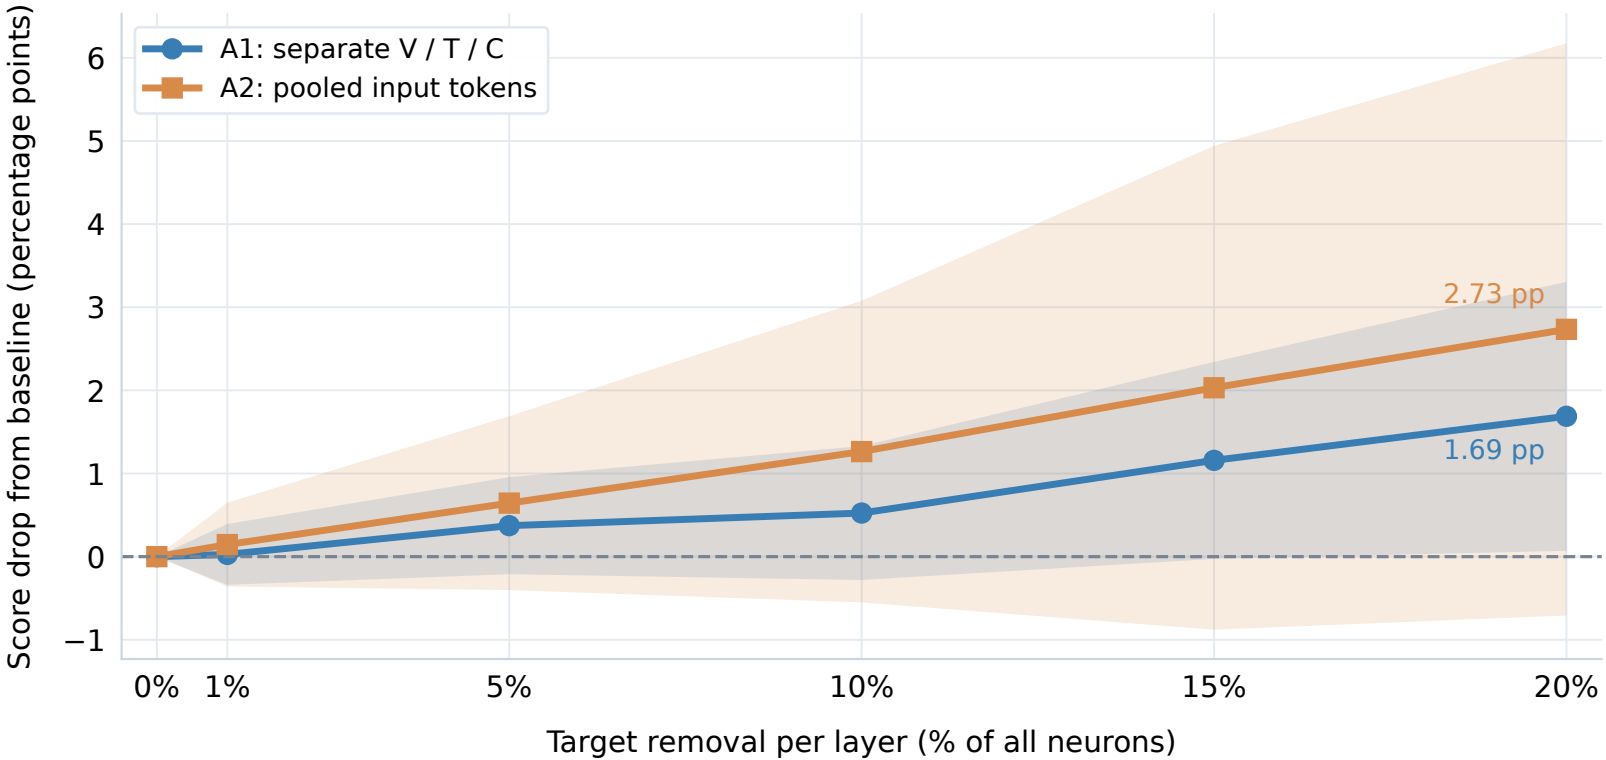

Shaded bands: ±1 SD across 16 model–dataset runs; descriptive spread, not confidence intervals.

Positive drop = worse performance. Different datasets use their established scoring metrics.